



Apache Spark* SQL expression evaluation

Chenzhao Guo

5/13/2017

Agenda

- Expression, UDF, UDAF & UDTF
- API definition
- Optimization in planning stage
- Code generation: part of Project Tungsten*
- Hivemall*: ML tools based on Hive UDF/UDAF/UDTF
- Future plan: vectorization

Agenda

- Expression, UDF, UDAF & UDTF
- API definition
- Optimization in planning stage
- Code generation: part of Project Tungsten*
- Hivemall*: ML tools based on Hive UDF/UDAF/UDTF
- Future plan: vectorization

Expression, UDF, UDAF & UDTF

- Expression(built-in function) represents a part of a structured query
 - Add(+), Coalesce, EqualTo(=), GreaterThan(>), Lower, Factorial, CurrentTimestamp(now)
 - Average, Count, Max (aggregate expression)

```
SELECT average(age) FROM test WHERE company='intel'
```

- Explode (generator)

```
SELECT explode(array(10, 20)) ->  
10  
20
```

- UDF(user defined function): a feature of Spark SQL* to define new column-based functions that extend the vocabulary of Spark SQL
 - UDAF(user defined aggregate function)
 - UDTF(user defined table function)

Agenda

- Expression, UDF, UDAF & UDTF
- **API definition**
- Optimization in planning stage
- Code generation: part of Project Tungsten*
- Hivemall*: ML tools based on Hive UDF/UDAF/UDTF
- Future plan: vectorization

API definition -- Expression

```
abstract class Expression {  
  def foldable: Boolean = false  
  def nullable: Boolean  
  def eval: Any  
  def genCode: ExprCode  
  ...  
}
```

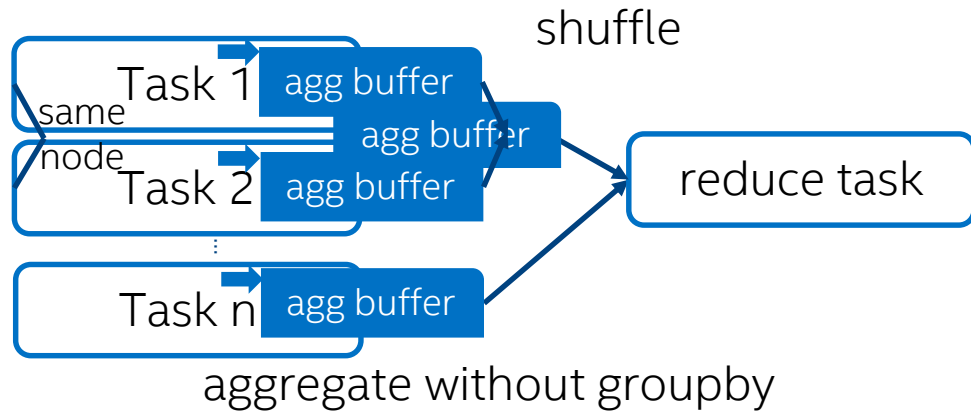
Add(1, 1)	Add(col1, col2)
true	false
false	col1.nullable col2.nullable

```
case class Add(left: Expression, right: Expression) extends BinaryArithmetic {  
  private lazy val numeric = TypeUtils.getNumeric(dataType)  
  protected override def nullSafeEval(input1: Any, input2: Any): Any = {  
    if (dataType.isInstanceOf[CalendarIntervalType]) {  
      input1.asInstanceOf[CalendarInterval].add(input2.asInstanceOf[CalendarInterval])  
    } else {  
      numeric.plus(input1, input2)  
    }  
  }  
  ...  
}
```

API definition -- Aggregate Expression



aggregation buffer for average



- Core API :
 - **Initialize** a new buffer
 - **Update** the buffer with a row
 - **Merge** 2 buffers together
 - **Evaluate** final result with the buffer

```
case class Average(child: Expression) extends DeclarativeAggregate {  
  override lazy val initialValues = Seq(  
    Literal(0),  
    Literal(0)  
  )  
  override lazy val updateExpressions = Seq(  
    Add(sum, child),  
    If(IsNull(child), count, count + 1L)  
  )  
  override lazy val mergeExpressions = Seq(  
    sum.left + sum.right,  
    count.left + count.right  
  )  
  override lazy val evaluateExpression = sum / count  
  ...  
}
```

API definition – Easier way to write & register UDF & UDAF

- Register UDF:

```
sqlContext.udf.register("strLen", (s: String) => s.length())
```

- Register UDAF:

```
class MyAggregate extends UserDefinedAggregateFunction {  
    ...  
}  
sqlContext.udf.register("my_aggregate", new MyAggregate )
```

- Pros: easier to implement compared to extending Expressions
- Cons: black box for Spark hence no optimization

Agenda

- Expression, UDF, UDAF & UDTF
- API definition
- **Optimization in planning stage**
- Code generation: part of Project Tungsten*
- Hivemall*: ML tools based on Hive UDF/UDAF/UDTF
- Future plan: vectorization

Optimization in planning stage

- Rules in analysis stage: PromoteStrings, ImplicitTypeCasts, DateTimeOperations, ...
- Rules in optimization stage: ConstantFolding, PushDownPredicate, LikeSimplification, ...

```
case class Concat(left: Expression, right: Expression) {  
  override def inputTypes: Seq[AbstractDataType] = Seq.fill(children.size)(StringType)  
  ...  
}
```

- concat example: *SELECT concat("Spark", "SQL") => "SparkSQL"*
 - SELECT concat(framework, version) from version_statistics*
- | id | framework | version |
|----|-----------|---------|
| 1 | "Spark" | 2.1 |
| 2 | "Flink" | 1.2 |
| 3 | "Hadoop" | 2.7 |

```
def implicitCast(e: Expression, expectedType: AbstractDataType): Option[Expression] = {  
  val ret: Expression = (e.datatype, expectedType) match {  
    case (any: AtomicType, StringType) if any != StringType => Cast(e, StringType)  
    ...  
  }  
}
```

Optimization in planning stage

- Constant Folding in optimization stage

```
abstract class Expression {  
  def foldable: Boolean = false  
  ...  
}
```

```
case class Add(left: Expression, right: Expression) {  
  override def foldable: Boolean = left.foldable && right.foldable  
  ...  
}
```

- Add(1,1) -> Add(Literal(1), Literal(1))
- Literal(1).foldable = true
- Add(1,1).foldable = true

➔ Add(1,1) -> 2

```
object ConstantFolding extends Rule[LogicalPlan] {  
  def apply(plan: LogicalPlan): LogicalPlan = plan transform {  
    case q: LogicalPlan => q transformExpressionsDown {  
      case l: Literal => l  
      case e if e.foldable => Literal.create(e.eval(EmptyRow), e.dataType)  
    }  
  }  
  ...  
}
```

Agenda


- Expression, UDF, UDAF & UDTF
- API definition
- Optimization in planning stage
- **Code generation: part of Project Tungsten***
- Hivemall*: ML tools based on Hive UDF/UDAF/UDTF
- Future plan: vectorization

Code generation: part of Project Tungsten*

- Code generation:
 - Collapse multiple expressions into one to eliminate overhead

```
abstract class Expression {  
  def genCode: ExprCode  
  ...  
}
```

- Project Tungsten*:
 - Hardware offers increasingly large network & disk IO bandwidth
 - Main focus: better efficiency of memory & CPU
 - *Also includes memory management & cache-aware computation besides codegen*



Evaluate col1+col2+col3 from leaf to root if without codegen

Agenda

- Expression, UDF, UDAF & UDTF
- API definition
- Optimization in planning stage
- Code generation: part of Project Tungsten*
- **Hivemall*: ML tools based on Hive UDF/UDAF/UDTF**
- Future plan: vectorization

Hivemall*: ML tools based on Hive UDF/UDAF/UDTF

- Machine learning on SQL
- Supports Hive, Spark, Pig
- Apache Incubator project

```
CREATE TABLE lr_model AS
SELECT      Reducers perform model averaging
            feature. in parallel
            avg(weight) as weight
FROM (
  SELECT logress(features, label, ..)
         as (feature, weight)
  FROM train
) t map-only task to learn a prediction model
GROUP BY feature
```

Image from <https://github.com/apache/incubator-hivemall>

- Other state-of-the-art machine learning algorithms: Soft Confidence Weighted, Adaptive Regularization of Weight Vectors, Factorization Machines, and AdaDelta

Agenda

- Expression, UDF, UDAF & UDTF
- API definition
- Optimization in planning stage
- Code generation: part of Project Tungsten*
- Hivemall*: ML tools based on Hive UDF/UDAF/UDTF
- **Future plan: vectorization**

- Future plan: vectorization
- Now: process data one row at a time
- One possible future: batch multiple rows together in a columnar format
- Already implemented: vectorized Parquet reader that does decompression and decoding in column batches

Row Format

1	john	4.1
2	mike	3.5
3	sally	6.4

Column Format

1	2	3
john	mike	sally
4.1	3.5	6.4



Benchmark from Databricks *



- Q & A

Legal Disclaimer

- No license (express or implied, by estoppel or otherwise) to any intellectual property rights is granted by this document.
- The products and services described may contain defects or errors known as errata which may cause deviations from published specifications. Current characterized errata are available on request.
- Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. No computer system can be absolutely secure. Check with your system manufacturer or retailer or learn more at [intel.com].
- Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.
- Intel, the Intel logo are trademarks of Intel Corporation in the U.S. and/or other countries.
- *Other names and brands may be claimed as the property of others.